

Deep Reinforcement Learning

John Schulman

OpenAI Berkeley¹

MLSS, May 2016, Cadiz

¹Berkeley Artificial Intelligence Research Lab

Agenda

Introduction and Overview

Markov Decision Processes

Reinforcement Learning via Black-Box Optimization

Policy Gradient Methods

Variance Reduction for Policy Gradients

Trust Region and Natural Gradient Methods


Open Problems

Course materials: goo.gl/5ws.gbJ

Introduction and Overview

What is Reinforcement Learning?

- ▶ Branch of machine learning concerned with taking sequences of actions
- ▶ Usually described in terms of agent interacting with a previously unknown environment, trying to maximize cumulative reward



Motor Control and Robotics



Robotics:

- ▶ Observations: camera images, joint angles
- ▶ Actions: joint torques
- ▶ Rewards: stay balanced, navigate to target locations, serve and protect humans

Business Operations

- ▶ Inventory Management
 - ▶ Observations: current inventory levels
 - ▶ Actions: number of units of each item to purchase
 - ▶ Rewards: profit
- ▶ Resource allocation: who to provide customer service to first
- ▶ Routing problems: in management of shipping fleet, which trucks / truckers to assign to which cargo

Games

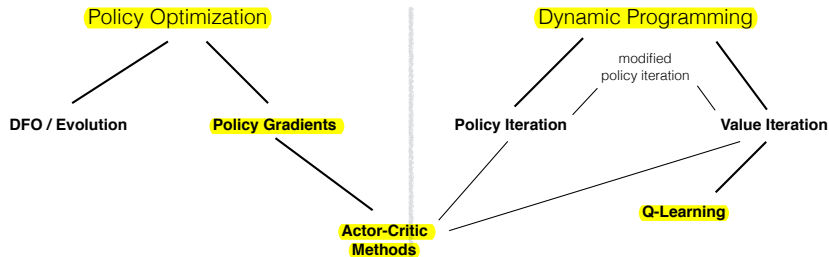
A different kind of optimization problem (min-max) but still considered to be RL.

- ▶ Go (complete information, deterministic) – *AlphaGo*²
- ▶ Backgammon (complete information, stochastic) – *TD-Gammon*³
- ▶ Stratego (incomplete information, deterministic)
- ▶ Poker (incomplete information, stochastic)

²David Silver, Aja Huang, et al. “Mastering the game of Go with deep neural networks and tree search”. In: *Nature* 529.7587 (2016), pp. 484–489.

³Gerald Tesauro. “Temporal difference learning and TD-Gammon”. In: *Communications of the ACM* 38.3 (1995), pp. 58–68.

Approaches to RL




What is Deep RL?

- ▶ RL using nonlinear function approximators
- ▶ Usually, updating parameters with stochastic gradient descent

What's Deep RL?

Whatever the front half of the cerebral cortex does (motor and executive cortices)



Markov Decision Processes

Definition

- ▶ Markov Decision Process (MDP) defined by $(\mathcal{S}, \mathcal{A}, P)$, where
 - ▶ \mathcal{S} : **state space**
 - ▶ \mathcal{A} : **action space**
 - ▶ $P(r, s' | s, a)$: a transition probability distribution
- ▶ Extra objects defined depending on problem setting
 - ▶ μ : Initial state distribution
 - ▶ γ : discount factor

Episodic Setting

- ▶ In each episode, the initial state is sampled from μ , and the process proceeds until the *terminal state* is reached.
For example:
 - ▶ Taxi robot reaches its destination (termination = good)
 - ▶ Waiter robot finishes a shift (fixed time)
 - ▶ Walking robot falls over (termination = bad)
- ▶ Goal: maximize expected reward per episode

Policies

- ▶ Deterministic policies: $a = \pi(s)$
- ▶ Stochastic policies: $a \sim \pi(a | s)$
- ▶ Parameterized policies: π_θ it is a function

Episodic Setting

$$s_0 \sim \mu(s_0)$$

$$a_0 \sim \pi(a_0 \mid s_0)$$

$$s_1, r_0 \sim P(s_1, r_0 \mid s_0, a_0)$$

$$a_1 \sim \pi(a_1 \mid s_1)$$

$$s_2, r_1 \sim P(s_2, r_1 \mid s_1, a_1)$$

...

$$a_{T-1} \sim \pi(a_{T-1} \mid s_{T-1})$$

$$s_T, r_{T-1} \sim P(s_T \mid s_{T-1}, a_{T-1})$$


Objective:

given a policy, the average reward (current + future)

maximize $\eta(\pi)$, where

$$\eta(\pi) = E[r_0 + r_1 + \cdots + r_{T-1} \mid \pi]$$

Episodic Setting



Objective:

maximize $\eta(\pi)$, where

$$\eta(\pi) = E[r_0 + r_1 + \dots + r_{T-1} \mid \pi]$$

Parameterized Policies

- ▶ A family of policies indexed by parameter vector $\theta \in \mathbb{R}^d$
 - ▶ Deterministic: $a = \pi(s, \theta)$
 - ▶ Stochastic: $\pi(a \mid s, \theta)$
- ▶ Analogous to classification or regression with input s , output a . E.g. for neural network stochastic policies:
 - ▶ Discrete action space: network outputs vector of probabilities
 - ▶ Continuous action space: network outputs mean and diagonal covariance of Gaussian

Reinforcement Learning via Black-Box Optimization

Derivative Free Optimization Approach

- ▶ Objective:

Maximize future rewards

$$\text{maximize } E[R \mid \pi(\cdot, \theta)]$$

- ▶ View $\theta \rightarrow \blacksquare \rightarrow R$ as a black box
- ▶ Ignore all other information other than R collected during episode

Cross-Entropy Method

- ▶ Evolutionary algorithm
- ▶ Works **embarrassingly** well

Method	Mean Score	Reference
Nonreinforcement learning		
Hand-coded	631,167	Dellacherie (Fahey, 2003)
Genetic algorithm	586,103	(Böhm et al., 2004)
Reinforcement learning		
Relational reinforcement learning+kernel-based regression	≈50	Ramon and Driessens (2004)
Policy iteration	3183	Bertsekas and Tsitsiklis (1996)
Least squares policy iteration	<3000	Lagoudakis, Parr, and Littman (2002)
Linear programming + Bootstrap	4274	Farias and van Roy (2006)
Natural policy gradient	≈6800	Kakade (2001)
CE+RL	21,252	
CE+RL, constant noise	72,705	
CE+RL, decreasing noise	348,895	

Approximate Dynamic Programming Finally Performs Well in the Game of Tetris

Victor Gabillon
INRIA Lille - Nord Europe,
Team Sequel, FRANCE
victor.gabillon@inria.fr

Mohammad Ghavamzadeh*
INRIA Lille - Team Sequel
& Adobe Research
mohammad.ghavamzadeh@inria.fr

Bruno Scherrer
INRIA Nancy - Grand Est,
Team Maia, FRANCE
bruno.scherrer@inria.fr

István Szita and András Lörincz. "Learning Tetris using the noisy cross-entropy method". In: *Neural computation* 18.12 (2006), pp. 2936–2941

Victor Gabillon, Mohammad Ghavamzadeh, and Bruno Scherrer. "Approximate Dynamic Programming Finally Performs Well in the Game of Tetris". In: *Advances in Neural Information Processing Systems*. 2013

Cross-Entropy Method

- ▶ Evolutionary algorithm
- ▶ Works **embarrassingly** well
- ▶ A similar algorithm, **Covariance Matrix Adaptation**, has become **standard in graphics**:

Optimal Gait and Form for Animal Locomotion

Kevin Wampler*

Zoran Popović

University of Washington




Optimizing Walking Controllers for Uncertain Inputs and Environments

Jack M. Wang

David J. Fleet

Aaron Hertzmann

University of Toronto



Cross-Entropy Method

Initialize $\mu \in \mathbb{R}^d, \sigma \in \mathbb{R}^d$

for iteration = 1, 2, ... **do**

Collect n samples of $\theta_i \sim N(\mu, \text{diag}(\sigma))$

Perform a noisy evaluation $R_i \sim \theta_i$

Select the top $p\%$ of samples (e.g. $p = 20$), which we'll
call the **elite set**

Fit a Gaussian distribution, with diagonal covariance,
to the elite set, obtaining a new μ, σ .

end for

Return the final μ .

Cross-Entropy Method

- ▶ Analysis: a very similar algorithm is an minorization-maximization (MM) algorithm, guaranteed to monotonically increase expected reward
- ▶ Recall that Monte-Carlo EM algorithm collects samples, reweights them, and then maximizes their logprob
- ▶ We can derive MM algorithm where each iteration you maximize $\sum_i \log p(\theta_i) R_i$

Policy Gradient Methods

Policy Gradient Methods: Overview

Problem:

Trajectory ?

$$\text{maximize } E[R \mid \pi_\theta]$$

Intuitions: collect a bunch of trajectories, and ...

1. Make the good trajectories more probable
2. Make the good actions more probable (actor-critic, GAE)
3. Push the actions towards good actions (DPG, SVG)

Score Function Gradient Estimator

- Consider an expectation $E_{x \sim p(x | \theta)}[f(x)]$. Want to compute gradient wrt θ

$$\begin{aligned}\nabla_{\theta} E_x[f(x)] &= \nabla_{\theta} \int dx \, p(x | \theta) f(x) \\ &= \int dx \, \nabla_{\theta} p(x | \theta) f(x) \\ &= \int dx \, p(x | \theta) \frac{\nabla_{\theta} p(x | \theta)}{p(x | \theta)} f(x) \\ &= \int dx \, p(x | \theta) \nabla_{\theta} \log p(x | \theta) f(x) \\ &= E_x[f(x) \nabla_{\theta} \log p(x | \theta)].\end{aligned}$$

- Last expression gives us **an unbiased gradient estimator**. Just sample $x_i \sim p(x | \theta)$, and compute $\hat{g}_i = f(x_i) \nabla_{\theta} \log p(x_i | \theta)$.
- Need to be able to compute and differentiate density $p(x | \theta)$ wrt θ

Derivation via Importance Sampling

Alternate Derivation Using Importance Sampling

$$\mathbb{E}_{x \sim \theta} [f(x)] = \mathbb{E}_{x \sim \theta_{\text{old}}} \left[\frac{p(x | \theta)}{p(x | \theta_{\text{old}})} f(x) \right]$$

$$\nabla_{\theta} \mathbb{E}_{x \sim \theta} [f(x)] = \mathbb{E}_{x \sim \theta_{\text{old}}} \left[\frac{\nabla_{\theta} p(x | \theta)}{p(x | \theta_{\text{old}})} f(x) \right]$$

$$\begin{aligned} \nabla_{\theta} \mathbb{E}_{x \sim \theta} [f(x)] \big|_{\theta=\theta_{\text{old}}} &= \mathbb{E}_{x \sim \theta_{\text{old}}} \left[\frac{\nabla_{\theta} p(x | \theta) \big|_{\theta=\theta_{\text{old}}}}{p(x | \theta_{\text{old}})} f(x) \right] \\ &= \mathbb{E}_{x \sim \theta_{\text{old}}} \left[\nabla_{\theta} \log p(x | \theta) \big|_{\theta=\theta_{\text{old}}} f(x) \right] \end{aligned}$$

Score Function Gradient Estimator: Intuition


$$\hat{g}_i = f(x_i) \nabla_{\theta} \log p(x_i | \theta)$$

- ▶ Let's say that $f(x)$ measures how good the sample x is.
- ▶ Moving in the direction \hat{g}_i pushes up the logprob of the sample, in proportion to how good it is
- ▶ *Valid even if $f(x)$ is discontinuous, and unknown, or sample space (containing x) is a discrete set*




Score Function Gradient Estimator: Intuition

$$\hat{g}_i = f(x_i) \nabla_{\theta} \log p(x_i | \theta)$$



Score Function Gradient Estimator: Intuition

$$\hat{g}_i = f(x_i) \nabla_{\theta} \log p(x_i | \theta)$$



Score Function Gradient Estimator for Policies

- ▶ Now random variable x is a whole trajectory

$$\tau = (s_0, a_0, r_0, s_1, a_1, r_1, \dots, s_{T-1}, a_{T-1}, r_{T-1}, s_T)$$

$$\nabla_{\theta} E_{\tau}[R(\tau)] = E_{\tau}[\nabla_{\theta} \log p(\tau | \theta) R(\tau)]$$

- ▶ Just need to write out $p(\tau | \theta)$:

$$p(\tau | \theta) = \mu(s_0) \prod_{t=0}^{T-1} [\pi(a_t | s_t, \theta) P(s_{t+1}, r_t | s_t, a_t)]$$

$$\log p(\tau | \theta) = \log \mu(s_0) + \sum_{t=0}^{T-1} [\log \pi(a_t | s_t, \theta) + \log P(s_{t+1}, r_t | s_t, a_t)]$$

$$\nabla_{\theta} \log p(\tau | \theta) = \nabla_{\theta} \sum_{t=0}^{T-1} \log \pi(a_t | s_t, \theta)$$

$$\nabla_{\theta} \mathbb{E}_{\tau}[R] = \mathbb{E}_{\tau} \left[R \nabla_{\theta} \sum_{t=0}^{T-1} \log \pi(a_t | s_t, \theta) \right]$$

- ▶ Interpretation: using good trajectories (high R) as supervised examples in classification / regression

Policy Gradient–Slightly Better Formula

- ▶ Previous slide:

$$\nabla_{\theta} \mathbb{E}_{\tau} [R] = \mathbb{E}_{\tau} \left[\left(\sum_{t=0}^{T-1} r_t \right) \left(\sum_{t=0}^{T-1} \nabla_{\theta} \log \pi(a_t | s_t, \theta) \right) \right]$$

- ▶ But we can cut trajectory to t steps and derive gradient estimator for one reward term $r_{t'}$.

$$\nabla_{\theta} \mathbb{E} [r_{t'}] = \mathbb{E} \left[r_{t'} \sum_{t=0}^t \nabla_{\theta} \log \pi(a_t | s_t, \theta) \right]$$

- ▶ Sum this formula over t , obtaining

$$\begin{aligned} \nabla_{\theta} \mathbb{E} [R] &= \mathbb{E} \left[\sum_{t=0}^{T-1} r_{t'} \sum_{t=0}^{t'} \nabla_{\theta} \log \pi(a_t | s_t, \theta) \right] \\ &= \mathbb{E} \left[\sum_{t=0}^{T-1} \nabla_{\theta} \log \pi(a_t | s_t, \theta) \sum_{t'=t}^{T-1} r_{t'} \right] \end{aligned}$$

Adding a Baseline

- ▶ Suppose $f(x) \geq 0, \quad \forall x$
- ▶ Then for every x_i , gradient estimator \hat{g}_i tries to push up it's density
- ▶ We can derive a new unbiased estimator that avoids this problem, and **only pushes up the density for better-than-average x_i .**

$$\begin{aligned}\nabla_{\theta} \mathbb{E}_x [f(x)] &= \nabla_{\theta} \mathbb{E}_x [f(x) - b] \\ &= \mathbb{E}_x [\nabla_{\theta} \log p(x | \theta) (f(x) - b)]\end{aligned}$$

- ▶ **A near-optimal choice of b is always $\mathbb{E}[f(x)]$ (which must be estimated)**

Policy Gradient with Baseline

- Recall

$$\nabla_{\theta} \mathbb{E}_{\tau} [R] = \sum_{t'=0}^{T-1} r_{t'} \sum_{t=t'}^{T-1} \nabla_{\theta} \log \pi(a_t | s_t, \theta)$$

- Using the $\mathbb{E}_{a_t} [\nabla_{\theta} \log \pi(a_t | s_t, \theta)] = 0$, we can show

$$\nabla_{\theta} \mathbb{E}_{\tau} [R] = \mathbb{E}_{\tau} \left[\sum_{t=0}^{T-1} \nabla_{\theta} \log \pi(a_t | s_t, \theta) \left(\sum_{t'=t}^{T-1} r_{t'} - b(s_t) \right) \right]$$

for any “baseline” function $b : \mathcal{S} \rightarrow \mathbb{R}$

- Increase logprob of action a_t proportionally to how much returns $\sum_{t=t'}^{T-1} r_{t'}$ are better than expected
- Later: use *value functions* to further isolate effect of action, at the cost of bias
- For more general picture of score function gradient estimator, see *stochastic computation graphs*⁴.

⁴ John Schulman, Nicolas Heess, et al. “Gradient Estimation Using Stochastic Computation Graphs”. In: *Advances in Neural Information Processing Systems*. 2015, pp. 3510–3522.

That's all for today

Course Materials: goo.gl/5wsGbJ

Variance Reduction for Policy Gradients

Review (I)

- Process for generating trajectory

$$\tau = (s_0, a_0, r_0, s_1, a_1, r_1, \dots, s_{T-1}, a_{T-1}, r_{T-1}, s_T)$$

$$s_0 \sim \mu(s_0)$$

$$a_0 \sim \pi(a_0 \mid s_0)$$

$$s_1, r_0 \sim P(s_1, r_0 \mid s_0, a_0)$$

$$a_1 \sim \pi(a_1 \mid s_1)$$

$$s_2, r_1 \sim P(s_2, r_1 \mid s_1, a_1)$$

...

$$a_{T-1} \sim \pi(a_{T-1} \mid s_{T-1})$$

$$s_T, r_{T-1} \sim P(s_T \mid s_{T-1}, a_{T-1})$$

- Given parameterized policy $\pi(a \mid s, \theta)$, the optimization problem is

$$\underset{\theta}{\text{maximize}} \mathbb{E}_{\tau} [R \mid \pi(\cdot \mid \cdot, \theta)]$$

where $R = r_0 + r_1 + \dots + r_{T-1}$.

Review (II)

- ▶ In general, we can compute gradients of expectations with the *score function gradient estimator*

$$\nabla_{\theta} \mathbb{E}_{x \sim p(x | \theta)} [f(x)] = \mathbb{E}_x [\nabla_{\theta} \log p(x | \theta) f(x)]$$

- ▶ We derived a formula for the policy gradient

$$\nabla_{\theta} \mathbb{E}_{\tau} [R] = \mathbb{E}_{\tau} \left[\sum_{t=0}^{T-1} \nabla_{\theta} \log \pi(a_t | s_t, \theta) \left(\sum_{t'=t}^{T-1} r_{t'} - b(s_t) \right) \right]$$

Value Functions

- ▶ The *state-value function* V^π is defined as:

$$V^\pi(s) = E[r_0 + r_1 + r_2 + \dots \mid s_0 = s]$$

Measures *expected future return, starting with state s*

- ▶ The *state-action value function* Q^π is defined as

$$Q^\pi(s, a) = E[r_0 + r_1 + r_2 + \dots \mid s_0 = s, a_0 = a]$$

- ▶ The *advantage function* A^π is

$$A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$$

Measures *how much better is action a than what the policy π would've done.*

Refining the Policy Gradient Formula

- Recall

$$\begin{aligned}\nabla_{\theta} \mathbb{E}_{\tau} [R] &= \mathbb{E}_{\tau} \left[\sum_{t=0}^{T-1} \nabla_{\theta} \log \pi(a_t | s_t, \theta) \left(\sum_{t'=t}^{T-1} r_{t'} - b(s_t) \right) \right] \\&= \sum_{t=0}^{T-1} \mathbb{E}_{\tau} \left[\nabla_{\theta} \log \pi(a_t | s_t, \theta) \left(\sum_{t'=t}^{T-1} r_{t'} - b(s_t) \right) \right] \\&= \sum_{t=0}^{T-1} \mathbb{E}_{s_0 \dots a_t} \left[\nabla_{\theta} \log \pi(a_t | s_t, \theta) \mathbb{E}_{r_t s_{t+1} \dots s_T} \left[\left(\sum_{t'=t}^{T-1} r_{t'} - b(s_t) \right) \right] \right] \\&= \sum_{t=0}^{T-1} \mathbb{E}_{s_0 \dots a_t} \left[\nabla_{\theta} \log \pi(a_t | s_t, \theta) \mathbb{E}_{r_t s_{t+1} \dots s_T} [Q^{\pi}(s_t, a_t) - b(s_t)] \right]\end{aligned}$$

- Where the last equality used the fact that

$$\mathbb{E}_{r_t s_{t+1} \dots s_T} \left[\sum_{t'=t}^{T-1} r_{t'} \right] = Q^{\pi}(s_t, a_t)$$

Refining the Policy Gradient Formula

- ▶ From the previous slide, we've obtained

$$\nabla_{\theta} \mathbb{E}_{\tau} [R] = \mathbb{E}_{\tau} \left[\sum_{t=0}^{T-1} \nabla_{\theta} \log \pi(a_t | s_t, \theta) (Q^{\pi}(s_t, a_t) - b(s_t)) \right]$$

- ▶ Now let's define $b(s) = V^{\pi}(s)$, which turns out to be near-optimal⁵. We get

$$\nabla_{\theta} \mathbb{E}_{\tau} [R] = \mathbb{E}_{\tau} \left[\sum_{t=0}^{T-1} \nabla_{\theta} \log \pi(a_t | s_t, \theta) A^{\pi}(s_t, a_t) \right]$$

- ▶ Intuition: increase the probability of good actions
(positive advantage) decrease the probability of bad ones
(negative advantage)

⁵Evan Greensmith, Peter L Bartlett, and Jonathan Baxter. "Variance reduction techniques for gradient estimates in reinforcement learning". In: *The Journal of Machine Learning Research* 5 (2004), pp. 1471–1530. 

Variance Reduction

- ▶ Now, we have the following policy gradient formula:

$$\nabla_{\theta} \mathbb{E}_{\tau} [R] = \mathbb{E}_{\tau} \left[\sum_{t=0}^{T-1} \nabla_{\theta} \log \pi(a_t | s_t, \theta) A^{\pi}(s_t, a_t) \right]$$


- ▶ A^{π} is not known, but we can plug in a random variable \hat{A}_t , an *advantage estimator*
- ▶ Previously, we showed that taking

$$\hat{A}_t = r_t + r_{t+1} + r_{t+2} + \dots - b(s_t)$$

for any function $b(s_t)$, gives an unbiased policy gradient estimator. $b(s_t) \approx V^{\pi}(s_t)$ gives variance reduction.

The Delayed Reward Problem

- ▶ One reason RL is difficult is the long delay between action and reward



The Delayed Reward Problem

- ▶ With policy gradient methods, we are confounding the effect of multiple actions:

$$\hat{A}_t = r_t + r_{t+1} + r_{t+2} + \dots - b(s_t)$$

mixes effect of $a_t, a_{t+1}, a_{t+2}, \dots$

- ▶ SNR of \hat{A}_t scales roughly as $1/T$
 - ▶ Only a_t contributes to *signal* $A^\pi(s_t, a_t)$, but a_{t+1}, a_{t+2}, \dots contribute to noise.

Var. Red. Idea 1: Using Discounts

- ▶ Discount factor γ , $0 < \gamma < 1$, downweights the effect of rewards that are far in the future—ignore long term dependencies
- ▶ We can form an advantage estimator using the *discounted return*:

$$\hat{A}_t^\gamma = \underbrace{r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots}_{\text{discounted return}} - b(s_t)$$

reduces to our previous estimator when $\gamma = 1$.

- ▶ So advantage has expectation zero, we should fit baseline to be *discounted value function*

$$V^{\pi, \gamma}(s) = \mathbb{E}_\tau [r_0 + \gamma r_1 + \gamma^2 r_2 + \dots \mid s_0 = s]$$

- ▶ \hat{A}_t^γ is a biased estimator of the advantage function

Var. Red. Idea 2: Value Functions in the Future

- ▶ Another approach for variance reduction is to use the value function to estimate future rewards

$$r_t + r_{t+1} + r_{t+2} + \dots \quad \text{use empirical rewards}$$

\Rightarrow

$$r_t + V(s_{t+1}) \quad \text{cut off at one timestep}$$

$$r_t + r_{t+1} + V(s_{t+2}) \quad \text{cut off at two timesteps}$$

\dots

Adding the baseline again, we get the advantage estimators

$$\hat{A}_t = r_t + V(s_{t+1}) - V(s_t) \quad \text{cut off at one timestep}$$

$$\hat{A}_t = r_t + r_{t+1} + V(s_{t+2}) - V(s_t) \quad \text{cut off at two timesteps}$$

\dots

Combining Ideas 1 and 2

- ▶ Can combine discounts and value functions in the future, e.g., $\hat{A}_t = r_t + \gamma V(s_{t+1}) - V(s_t)$, where V approximates discounted value function $V^{\pi, \gamma}$.
- ▶ The above formula is called an *actor-critic* method, where *actor* is the policy π , and *critic* is the value function V .⁶
- ▶ Going further, the *generalized advantage estimator*⁷

$$\hat{A}_t^{\gamma, \lambda} = \delta_t + (\gamma\lambda)\delta_{t+1} + (\gamma\lambda)^2\delta_{t+2} + \dots$$

$$\text{where } \delta_t = r_t + \gamma V(s_{t+1}) - V(s_t)$$

- ▶ Interpolates between two previous estimators:

$$\lambda = 0 : \quad r_t + \gamma V(s_{t+1}) - V(s_t) \quad (\text{low } v, \text{ high } b)$$

$$\lambda = 1 : \quad r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots - V(s_t) \quad (\text{low } b, \text{ high } v)$$

⁶Vijay R Konda and John N Tsitsiklis. "Actor-Critic Algorithms." In: *Advances in Neural Information Processing Systems*. Vol. 13. Citeseer. 1999, pp. 1008–1014.

⁷John Schulman, Philipp Moritz, et al. "High-dimensional continuous control using generalized advantage estimation". In: *arXiv preprint arXiv:1506.02438* (2015).

Alternative Approach: Reparameterization

- ▶ Suppose problem has continuous action space, $a \in \mathbb{R}^d$
- ▶ Then $\frac{d}{da} Q^\pi(s, a)$ tells use how to improve our action
- ▶ We can use reparameterization trick, so a is a deterministic function $a = f(s, z)$, where z is noise. Then,

$$\nabla_\theta \mathbb{E}_\tau [R] = \nabla_\theta Q^\pi(s_0, a_0) + \nabla_\theta Q^\pi(s_1, a_1) + \dots$$

- ▶ This method is called the deterministic policy gradient⁸
- ▶ A generalized version, which also uses a dynamics model, is described as the stochastic value gradient⁹

⁸David Silver, Guy Lever, et al. "Deterministic policy gradient algorithms". In: *ICML 2014*; Timothy P Lillicrap et al. "Continuous control with deep reinforcement learning". In: *arXiv preprint arXiv:1509.02971* (2015).

⁹Nicolas Heess et al. "Learning continuous control policies by stochastic value gradients". In: *Advances in Neural Information Processing Systems*. 2015, pp. 2926–2934.

Trust Region and Natural Gradient Methods

Optimization Issues with Policy Gradients

- ▶ Hard to choose reasonable stepsize that works for the whole optimization
 - ▶ we have a gradient estimate, no objective for line search
 - ▶ statistics of data (observations and rewards) change during learning
- ▶ They make inefficient use of data: each experience is only used to compute one gradient.
 - ▶ Given a batch of trajectories, what's the most we can do with it?

Policy Performance Function

- ▶ Let $\eta(\pi)$ denote the performance of policy π

$$\eta(\pi) = \mathbb{E}_{\tau} [R|\pi]$$

- ▶ The following neat identity holds:

$$\eta(\tilde{\pi}) = \eta(\pi) + \mathbb{E}_{\tau \sim \tilde{\pi}} [A^{\pi}(s_0, a_0) + A^{\pi}(s_1, a_1) + A^{\pi}(s_2, a_2) + \dots]$$

- ▶ Proof: consider nonstationary policy $\pi_0\pi_1\pi_2, \dots$

$$\begin{aligned} \eta(\tilde{\pi}\tilde{\pi}\tilde{\pi}\dots) &= \eta(\pi\pi\pi\dots) \\ &\quad + \eta(\tilde{\pi}\pi\pi\dots) - \eta(\pi\pi\pi\dots) \\ &\quad + \eta(\tilde{\pi}\tilde{\pi}\pi\dots) - \eta(\tilde{\pi}\pi\pi\dots) \\ &\quad + \eta(\tilde{\pi}\tilde{\pi}\tilde{\pi}\dots) - \eta(\tilde{\pi}\tilde{\pi}\pi\dots) \\ &\quad + \dots \end{aligned}$$

- ▶ t^{th} difference term equals $A^{\pi}(s_t, a_t)$

Local Approximation

- ▶ We just derived an expression for the performance of a policy $\tilde{\pi}$ relative to π

$$\begin{aligned}\eta(\tilde{\pi}) &= \eta(\pi) + \mathbb{E}_{\tau \sim \tilde{\pi}} [A^{\pi}(s_0, a_0) + A^{\pi}(s_1, a_1) + \dots] \\ &= \eta(\pi) + \mathbb{E}_{s_{0:\infty} \sim \tilde{\pi}} [\mathbb{E}_{a_{0:\infty} \sim \tilde{\pi}} [A^{\pi}(s_0, a_0) + A^{\pi}(s_1, a_1) + \dots]]\end{aligned}$$

- ▶ Can't use this to optimize $\tilde{\pi}$ because state distribution has complicated dependence.
- ▶ Let's define L_{π} the *local approximation*, which ignores change in state distribution—can be estimated by sampling from π

$$\begin{aligned}L_{\pi}(\tilde{\pi}) &= \mathbb{E}_{s_{0:\infty} \sim \pi} [\mathbb{E}_{a_{0:\infty} \sim \tilde{\pi}} [A^{\pi}(s_0, a_0) + A^{\pi}(s_1, a_1) + \dots]] \\ &= \mathbb{E}_{s_{0:\infty}} \left[\sum_{t=0}^{T-1} \mathbb{E}_{a \sim \tilde{\pi}} [A^{\pi}(s_t, a_t)] \right] \\ &= \mathbb{E}_{s_{0:\infty}} \left[\sum_{t=0}^{T-1} \mathbb{E}_{a \sim \pi} \left[\frac{\tilde{\pi}(a_t | s_t)}{\pi(a_t | s_t)} A^{\pi}(s_t, a_t) \right] \right] \\ &= \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^{T-1} \frac{\tilde{\pi}(a_t | s_t)}{\pi(a_t | s_t)} A^{\pi}(s_t, a_t) \right]\end{aligned}$$

Local Approximation

- ▶ Now let's consider parameterized policy, $\pi(a | s, \theta)$. Sample with θ_{old} , now write local approximation in terms of θ .

$$L_{\pi}(\tilde{\pi}) = \mathbb{E}_{s_{0:\infty}} \left[\sum_{t=0}^{T-1} \mathbb{E}_{a \sim \pi} \left[\frac{\tilde{\pi}(a_t | s_t)}{\pi(a_t | s_t)} A^{\pi}(s_t, a_t) \right] \right]$$
$$\Rightarrow L_{\theta_{\text{old}}}(\theta) = \mathbb{E}_{s_{0:\infty}} \left[\sum_{t=0}^{T-1} \mathbb{E}_{a \sim \theta} \left[\frac{\pi(a_t | s_t, \theta)}{\pi(a_t | s_t, \theta_{\text{old}})} A^{\theta}(s_t, a_t) \right] \right]$$


- ▶ $L_{\theta_{\text{old}}}(\theta)$ matches $\eta(\theta)$ to first order around θ_{old} .

$$\begin{aligned} \nabla_{\theta} L_{\theta_{\text{old}}}(\theta) \big|_{\theta=\theta_0} &= \mathbb{E}_{s_{0:\infty}} \left[\sum_{t=0}^{T-1} \mathbb{E}_{a \sim \theta} \left[\frac{\nabla_{\theta} \pi(a_t | s_t, \theta)}{\pi(a_t | s_t, \theta_{\text{old}})} A^{\theta}(s_t, a_t) \right] \right] \\ &= \mathbb{E}_{s_{0:\infty}} \left[\sum_{t=0}^{T-1} \mathbb{E}_{a \sim \theta} \left[\nabla_{\theta} \log \pi(a_t | s_t, \theta) A^{\theta}(s_t, a_t) \right] \right] \\ &= \nabla_{\theta} \eta(\theta) \big|_{\theta=\theta_{\text{old}}} \end{aligned}$$

MM Algorithm

- ▶ Theorem (ignoring some details)¹⁰

$$\eta(\theta) \geq \underbrace{L_{\theta_{\text{old}}}(\theta)}_{\text{local approx. to } \eta} - \underbrace{C \max_s D_{KL} [\pi(\cdot | \theta_{\text{old}}, s) \parallel \pi(\cdot | \theta, s)]}_{\text{penalty for changing policy}}$$



- ▶ If $\theta_{\text{old}} \rightarrow \theta_{\text{new}}$ improves lower bound, it's guaranteed to improve η

¹⁰ John Schulman, Sergey Levine, et al. "Trust Region Policy Optimization". In: *arXiv preprint arXiv:1502.05477* (2015).

Review

- ▶ Want to optimize $\eta(\theta)$. Collected data with policy parameter θ_{old} , now want to do update
- ▶ Derived local approximation $L_{\theta_{\text{old}}}(\theta)$
- ▶ Optimizing KL penalized local approximation gives guaranteed improvement to η
- ▶ More approximations gives practical algorithm, called TRPO

TRPO—Approximations

- ▶ Steps:
 - ▶ Instead of max over state space, take mean
 - ▶ Linear approximation to L , quadratic approximation to KL divergence
 - ▶ Use hard constraint on KL divergence instead of penalty
- ▶ Solve the following problem approximately


$$\begin{aligned} & \text{maximize } L_{\theta_{\text{old}}}(\theta) \\ & \text{subject to } \overline{D}_{KL}[\theta_{\text{old}} \parallel \theta] \leq \delta \end{aligned}$$

- ▶ Solve approximately through line search in the *natural gradient* direction $s = F^{-1}g$
- ▶ Resulting algorithm is a refined version of *natural policy gradient*¹¹


¹¹Sham Kakade. "A Natural Policy Gradient." In: *NIPS*. vol. 14. 2001, pp. 1531–1538.

Empirical Results: TRPO + GAE

- ▶ TRPO, with neural network policies, was applied to learn controllers for 2D robotic swimming, hopping, and walking, and playing Atari games¹²



- ▶ Used TRPO along with generalized advantage estimation to optimize locomotion policies for 3D simulated robots¹³



¹²John Schulman, Sergey Levine, et al. "Trust Region Policy Optimization". In: *arXiv preprint arXiv:1502.05477* (2015).

¹³John Schulman, Philipp Moritz, et al. "High-dimensional continuous control using generalized advantage estimation". In: *arXiv preprint arXiv:1506.02438* (2015).

Putting In Perspective

Quick and incomplete overview of recent results with deep RL algorithms

- ▶ Policy gradient methods
 - ▶ TRPO + GAE
 - ▶ Standard policy gradient (no trust region) + deep nets + parallel implementation¹⁴
 - ▶ Repair trick¹⁵
- ▶ Q-learning¹⁶ and modifications¹⁷
- ▶ Combining search + supervised learning¹⁸

¹⁴V. Mnih et al. "Playing Atari with Deep Reinforcement Learning". In: *arXiv preprint arXiv:1312.5602* (2013).

¹⁵Nicolas Heess et al. "Learning continuous control policies by stochastic value gradients". In: *Advances in Neural Information Processing Systems*. 2015, pp. 2926–2934; Timothy P Lillicrap et al. "Continuous control with deep reinforcement learning". In: *arXiv preprint arXiv:1509.02971* (2015).

¹⁶V. Mnih et al. "Playing Atari with Deep Reinforcement Learning". In: *arXiv preprint arXiv:1312.5602* (2013).

¹⁷Ziyu Wang, Nando de Freitas, and Marc Lanctot. "Dueling Network Architectures for Deep Reinforcement Learning". In: *arXiv preprint arXiv:1511.06581* (2015); Hado V Hasselt. "Double Q-learning". In: *Advances in Neural Information Processing Systems*. 2010, pp. 2613–2621.

¹⁸X. Guo et al. "Deep learning for real-time Atari game play using offline Monte-Carlo tree search planning". In: *Advances in Neural Information Processing Systems*. 2014, pp. 3338–3346; Sergey Levine et al. "End-to-end training of deep visuomotor policies". In: *arXiv preprint arXiv:1504.00702* (2015); Igor Mordatch et al. "Interactive Control of Diverse Complex Characters with Neural Networks". In: *Advances in Neural Information Processing Systems*. 2015, pp. 3114–3122.

Open Problems

What's the Right Core Model-Free Algorithm?

- ▶ Policy gradients (score function vs. reparameterization, natural vs. not natural) vs. Q-learning vs. derivative-free optimization vs others
- ▶ Desiderata
 - ▶ scalable
 - ▶ sample-efficient
 - ▶ robust
 - ▶ learns from *off-policy* data

Exploration


- ▶ Exploration: actively encourage agent to reach unfamiliar parts of state space, avoid getting stuck in local maximum of performance
- ▶ Can solve finite MDPs in polynomial time with exploration¹⁹
 - ▶ optimism about new states and actions
 - ▶ maintain distribution over possible models, and plan with them (Bayesian RL, Thompson sampling)
- ▶ How to do exploration in deep RL setting? Thompson sampling²⁰, novelty bonus²¹

¹⁹Alexander L Strehl et al. "PAC model-free reinforcement learning". In: *Proceedings of the 23rd international conference on Machine learning*. ACM. 2006, pp. 881–888.

²⁰Ian Osband et al. "Deep Exploration via Bootstrapped DQN". . In: *arXiv preprint arXiv:1602.04621* (2016).

²¹Bradly C Stadie, Sergey Levine, and Pieter Abbeel. "Incentivizing Exploration In Reinforcement Learning With Deep Predictive Models". In: *arXiv preprint arXiv:1507.00814* (2015).

Hierarchy



----->
walk to x ... fetch object y ... say z01 hz: 10^3 time steps per day

----->
footstep planning: 1hz: 10^5 timesteps / day

----->
torque control: 100hz: 10^7 timesteps / day

More Open Problems

- ▶ Using learned models
- ▶ Learning from demonstrations

The End

Questions?