

```
import random
import numpy as np
import tensorflow as tf
from replay_buffer import ReplayBuffer

class NeuralQLearner(object):

    def __init__(self, session,
                  optimizer,
                  q_network,
                  state_dim,
                  num_actions,
                  batch_size=32,
                  init_exp=0.5,      # initial exploration prob
                  final_exp=0.1,     # final exploration prob
                  anneal_steps=10000, # N steps for annealing exploration
                  replay_buffer_size=10000,
                  store_replay_every=5, # how frequent to store experience
                  discount_factor=0.9, # discount future rewards
                  target_update_rate=0.01,
                  reg_param=0.01, # regularization constants
                  max_gradient=5, # max gradient norms
                  double_q_learning=False,
                  summary_writer=None,
                  summary_every=100):

        # tensorflow machinery
        self.session = session
        self.optimizer = optimizer
        self.summary_writer = summary_writer

        # model components
        self.q_network = q_network
        self.replay_buffer = ReplayBuffer(buffer_size=replay_buffer_size)

        # Q learning parameters
        self.batch_size = batch_size
        self.state_dim = state_dim
        self.num_actions = num_actions
        self.exploration = init_exp
        self.init_exp = init_exp
        self.final_exp = final_exp
        self.anneal_steps = anneal_steps
        self.discount_factor = discount_factor
        self.target_update_rate = target_update_rate
        self.double_q_learning = double_q_learning

        # training parameters
        self.max_gradient = max_gradient
        self.reg_param = reg_param

        # counters
        self.store_replay_every = store_replay_every
        self.store_experience_cnt = 0
        self.train_iteration = 0

        # create and initialize variables
        self.create_variables()
        var_lists = tf.get_collection(tf.GraphKeys.VARIABLES)
        self.session.run(tf.initialize_variables(var_lists))

        # make sure all variables are initialized
        self.session.run(tf.assert_variables_initialized())
```

```

if self.summary_writer is not None:
    # graph was not available when journalist was created
    self.summary_writer.add_graph(self.session.graph)
    self.summary_every = summary_every

def create_variables(self):
    # compute action from a state:  $a^* = \operatorname{argmax}_a Q(s_t, a)$ 
    with tf.name_scope("predict_actions"):
        # raw state representation
        self.states = tf.placeholder(tf.float32, (None, self.state_dim), name="states")
        # initialize Q network
        with tf.variable_scope("q_network"):
            self.q_outputs = self.q_network(self.states)
        # predict actions from Q network
        self.action_scores = tf.identity(self.q_outputs, name="action_scores")
        tf.histogram_summary("action_scores", self.action_scores)
        self.predicted_actions = tf.argmax(self.action_scores, dimension=1, name="predicted_actions")

    # estimate rewards using the next state:  $r(s_t, a_t) + \operatorname{argmax}_a Q(s_{t+1}, a)$ 
    with tf.name_scope("estimate_future_rewards"):
        self.next_states = tf.placeholder(tf.float32, (None, self.state_dim), name="next_states")
        self.next_state_mask = tf.placeholder(tf.float32, (None,), name="next_state_masks")

        if self.double_q_learning:
            # reuse Q network for action selection
            with tf.variable_scope("q_network", reuse=True):
                self.q_next_outputs = self.q_network(self.next_states)
            self.action_selection = tf.argmax(tf.stop_gradient(self.q_next_outputs), 1,
            name="action_selection")
            tf.histogram_summary("action_selection", self.action_selection)
            self.action_selection_mask = tf.one_hot(self.action_selection, self.num_actions, 1, 0)
            # use target network for action evaluation
            with tf.variable_scope("target_network"):
                self.target_outputs = self.q_network(self.next_states) *
tf.cast(self.action_selection_mask, tf.float32)
            self.action_evaluation = tf.reduce_sum(self.target_outputs, reduction_indices=[1,])
            tf.histogram_summary("action_evaluation", self.action_evaluation)
            self.target_values = self.action_evaluation * self.next_state_mask
        else:
            # initialize target network
            with tf.variable_scope("target_network"):
                self.target_outputs = self.q_network(self.next_states)
            # compute future rewards
            self.next_action_scores = tf.stop_gradient(self.target_outputs)
            self.target_values = tf.reduce_max(self.next_action_scores, reduction_indices=[1,]) *
self.next_state_mask
            tf.histogram_summary("next_action_scores", self.next_action_scores)

        self.rewards = tf.placeholder(tf.float32, (None,), name="rewards")
        self.future_rewards = self.rewards + self.discount_factor * self.target_values

    # compute loss and gradients
    with tf.name_scope("compute_temporal_differences"):
        # compute temporal difference loss
        self.action_mask = tf.placeholder(tf.float32, (None, self.num_actions), name="action_mask")
        self.masked_action_scores = tf.reduce_sum(self.action_scores * self.action_mask,
reduction_indices=[1,])
        self.temp_diff = self.masked_action_scores - self.future_rewards
        self.td_loss = tf.reduce_mean(tf.square(self.temp_diff))
        # regularization loss
        q_network_variables = tf.get_collection(tf.GraphKeys.TRAINABLE_VARIABLES, scope="q_network")
        self.reg_loss = self.reg_param * tf.reduce_sum([tf.reduce_sum(tf.square(x)) for x in
q_network_variables])
        # compute total loss and gradients
        self.loss = self.td_loss + self.reg_loss

```

```

gradients = self.optimizer.compute_gradients(self.loss)
# clip gradients by norm
for i, (grad, var) in enumerate(gradients):
    if grad is not None:
        gradients[i] = (tf.clip_by_norm(grad, self.max_gradient), var)
# add histograms for gradients.
for grad, var in gradients:
    tf.histogram_summary(var.name, var)
    if grad is not None:
        tf.histogram_summary(var.name + '/gradients', grad)
self.train_op = self.optimizer.apply_gradients(gradients)

# update target network with Q network
with tf.name_scope("update_target_network"):
    self.target_network_update = []
    # slowly update target network parameters with Q network parameters
    q_network_variables = tf.get_collection(tf.GraphKeys.TRAINABLE_VARIABLES, scope="q_network")
    target_network_variables = tf.get_collection(tf.GraphKeys.TRAINABLE_VARIABLES,
scope="target_network")
    for v_source, v_target in zip(q_network_variables, target_network_variables):
        # this is equivalent to target = (1-alpha) * target + alpha * source
        update_op = v_target.assign_sub(self.target_update_rate * (v_target - v_source))
        self.target_network_update.append(update_op)
    self.target_network_update = tf.group(*self.target_network_update)

# scalar summaries
tf.scalar_summary("td_loss", self.td_loss)
tf.scalar_summary("reg_loss", self.reg_loss)
tf.scalar_summary("total_loss", self.loss)
tf.scalar_summary("exploration", self.exploration)

self.summarize = tf.merge_all_summaries()
self.no_op = tf.no_op()

def storeExperience(self, state, action, reward, next_state, done):
    # always store end states
    if self.store_experience_cnt % self.store_replay_every == 0 or done:
        self.replay_buffer.add(state, action, reward, next_state, done)
    self.store_experience_cnt += 1

def eGreedyAction(self, states, explore=True):
    if explore and self.exploration > random.random():
        return random.randint(0, self.num_actions-1)
    else:
        return self.session.run(self.predicted_actions, {self.states: states})[0]

def annealExploration(self, strategy='linear'):
    ratio = max((self.anneal_steps - self.train_iteration)/float(self.anneal_steps), 0)
    self.exploration = (self.init_exp - self.final_exp) * ratio + self.final_exp

def updateModel(self):
    # not enough experiences yet
    if self.replay_buffer.count() < self.batch_size:
        return

    batch = self.replay_buffer.getBatch(self.batch_size)
    states = np.zeros((self.batch_size, self.state_dim))
    rewards = np.zeros((self.batch_size,))
    action_mask = np.zeros((self.batch_size, self.num_actions))
    next_states = np.zeros((self.batch_size, self.state_dim))
    next_state_mask = np.zeros((self.batch_size,))

    for k, (s0, a, r, s1, done) in enumerate(batch):
        states[k] = s0
        rewards[k] = r

```

```
    action_mask[k][a] = 1
    # check terminal state
    if not done:
        next_states[k] = s1
        next_state_mask[k] = 1

    # whether to calculate summaries
    calculate_summaries = self.train_iteration % self.summary_every == 0 and self.summary_writer is
not None

    # perform one update of training
    cost, _, summary_str = self.session.run([
        self.loss,
        self.train_op,
        self.summarize if calculate_summaries else self.no_op
    ], {
        self.states:          states,
        self.next_states:     next_states,
        self.next_state_mask: next_state_mask,
        self.action_mask:     action_mask,
        self.rewards:         rewards
    })

    # update target network using Q-network
    self.session.run(self.target_network_update)

    # emit summaries
    if calculate_summaries:
        self.summary_writer.add_summary(summary_str, self.train_iteration)

    self.annealExploration()
    self.train_iteration += 1
```